

The Case for Hedonic Primacy

A Framework for Consciousness, Adaptation,
and the Engineering of Experience

Thomas Peacock

University of Queensland

April 2026

Abstract

This paper sets out a philosophical framework I have arrived at through a combination of first-person experience with altered hedonic states, engagement with the neuroscience of reward and adaptation, and extended dialogue with the existing literature on ethics and personal identity. The framework rests on a single foundational claim: positive valence is the only thing that is intrinsically good, and negative valence is the only thing that is intrinsically bad. Every secondary value that humans pursue — meaning, purpose, dignity, accomplishment, freedom, authenticity, even love — derives its value from the hedonic experiences it produces, and nothing else. I argue that the apparent disagreement with this claim in both common intuition and academic philosophy is not evidence against it, but rather an artifact of two related problems: biological hedonic adaptation, which prevents sustained positive experience and therefore conditions us to treat wellbeing as a scarce luxury rather than the point of existence, and the phenomenological asymmetry between peak hedonic states and sober reflection, which means that the sober mind cannot fully access the evidence on which hedonic primacy rests. From the foundational claim I derive a consciousness-centric rather than species-centric ethics, an intensity-threshold account of suffering, a two-level structure distinguishing endpoint ideals from transitional rules, and an account of personal identity grounded in open individualism and gradual substrate replacement that permits cosmic-scale hedonic engineering without requiring death of the self. I close by stating the precautionary conditions under which that engineering would be permissible, and the conditions under which it would not. The framework is presented as revisable in principle. Its core, however, I hold to be as close to self-evidently true as any ethical claim has ever been, for reasons the paper tries to make visible to the reader who has not yet felt them.

1. Statement of the Problem

A note on predecessors before I begin. The position defended in this paper overlaps in its core commitments with David Pearce's *The Hedonistic Imperative* (1995) and with the research programme of Andrés Gómez Emilsson at the Qualia Research Institute. I did not arrive at these commitments by reading either of them; I found their work after independently reaching conclusions that turned out to match theirs on several important points. I acknowledge the overlap openly and I am indebted to both thinkers for having developed these ideas in print before I got to them. Where I believe I have added something — the phenomenological-asymmetry argument, the specific two-level structure, the precautionary-safeguard response to the cataclysmic-risk problem, and the particular formulation of open individualism as one subject exchanging information with itself across structurally-separated interfaces — I have tried to make the addition visible rather than smuggled. Where I am repeating what has already been said, I am repeating it deliberately, because I think the position has not been said loudly enough.

I want to start with something that seems obvious and will turn out not to be. Every conscious being that has ever existed on this planet has preferred pleasant experience to unpleasant experience. Not as a cultural accident, not as a contingent psychological fact, but as a structural feature of what it is to be conscious at all. A being that did not prefer pleasure to pain would not be recognisable as a conscious subject; it would be something we would describe as a process, a mechanism, a system, but not as someone to whom anything matters. The fact that something matters to a being is the same fact as that being having hedonic states. There is no other place the mattering can come from, because there is no other ingredient in consciousness that could do the job.

If this is right — and I will argue that it is, and that the arguments against it fail for specific reasons I can name — then the question of what we should do with our lives, our species, our technology, and our universe admits a much simpler answer than philosophy has historically been willing to accept. We should be trying to produce positive hedonic experience and prevent negative hedonic experience in every conscious being capable of having either. That is the whole project. Everything else is instrumental to it or a distraction from it.

The reason this answer has not been taken seriously is not that it is wrong. It is that it is impossible to live consistently with, given the biology we currently have. Our brains are designed by evolutionary selection to use hedonic states as a control signal for survival behaviour, not as a goal in themselves. The signal is transient by design: if positive valence persisted, it would stop functioning as a gradient to climb. Adaptation — the downregulation of reward circuitry in response to sustained input — is not a bug in human happiness. It is the architecture that makes human happiness impossible in the form that would actually solve the problem. Every person who has ever tried to build a life of sustained positive experience has hit this wall. The wall is not psychological, spiritual, or moral. It is receptor density, transcription factor accumulation, synaptic pruning, and efficient coding in the central nervous system. It is physics, and no amount of wisdom will unbuild it.

What I am going to argue in this paper is that once we see this clearly, the entire map of what ethics has been doing for the last three thousand years reorganises. Most of what philosophers have called virtue, meaning, purpose, and the good life turn out to be coping

strategies for a substrate that cannot deliver the thing it actually wants. And the goal of ethics, properly understood, is not to refine the coping strategies. It is to change the substrate.

2. The Foundational Claim

The framework begins with a foundational claim: positive valence is the only intrinsic good, and negative valence is the only intrinsic bad. Every secondary value, including meaning, purpose, dignity, accomplishment, and authenticity, derives its value from its relationship to subjective hedonic experience. Meaning is valued because its presence produces positive affect and its absence produces negative affect. Purpose is valued because purposelessness produces suffering. Dignity matters because its violation causes distress. This is not a reductive move that denies the importance of these things. It is a move that locates them correctly, so that we can actually understand why they matter and, more importantly, what else would matter if we could produce the same affective result without them.

The test for whether something is intrinsically good or instrumentally good is simple and I do not think it has ever been answered cleanly by the people who have tried to resist this conclusion. Imagine any candidate for intrinsic value — say, knowledge, or freedom, or justice. Now imagine a being that has that thing in full measure but experiences no positive affect from having it and no negative affect from its absence. Is that being better off for having the thing? If you answer yes, I want you to locate the improvement. Where is it? In whom? Whose experience is better? If you cannot point to a conscious being for whom the thing produces some affective difference, you have not described a state of the world that is better or worse than any other. You have described a rearrangement of atoms that nobody cares about. The improvement you thought you were tracking is a shadow cast by your own positive affect at the idea of the thing, projected onto the state of the world, and then confused for a property of the world itself.

This is sometimes called the experience requirement on value, and the standard objection to it is G. E. Moore's open-question argument: granting that everyone wants positive valence, we can still coherently ask whether getting what everyone wants is actually good. Moore held that the coherence of the question was evidence that wanting and goodness are not identical. I reject this for a specific reason. The coherence of the question is itself an artifact of the questioner being in a sober reflective state, evaluating from outside any particular hedonic experience. From inside intense positive valence — and I will have more to say about this in the section on phenomenological asymmetry — the question simply does not arise. It is not that it arises and is answered trivially; it is that the cognitive apparatus that would generate such a question is suspended. The coherence of Moore's question, in other words, is evidence about the state of the questioner, not about the relationship between wanting and goodness.

A second objection is the so-called experience machine thought experiment from Robert Nozick, which asks whether you would plug into a machine guaranteed to produce any experience you wanted, and notes that most people say no. I take the "no" to be diagnostic of exactly the problem the framework is trying to solve. People say no because, in their current substrate, plugging in would entail the loss of relationships, projects, and narrative identity,

all of which are themselves sources of current hedonic states that they would not want to forgo. The refusal is not evidence that something other than experience matters. It is evidence that people are attached to the specific shape of their current life because that shape is where their current positive states come from. Ask the question differently: if you could be guaranteed, with certainty, that your experience would become uniformly better than it is now, by any measure you currently use to decide whether you are having a good day, and that the decision was reversible, would you do it? Most people say yes. The difference between the two framings is that the first preserves the intuition that the current substrate is the baseline of reality and any deviation from it is a loss, while the second removes that framing effect. The framing effect is the whole objection. Without it, the experience machine ceases to be a counterexample and becomes an intuition pump for precisely the position it was designed to refute.

3. The Adaptation Problem and Why Biological Happiness is Impossible

The central empirical premise of this paper, and the one on which the entire moral case for hedonic engineering rests, is that biological consciousness cannot sustain high positive valence. This is not a claim about psychology, resilience, or attitude. It is a claim about the physical architecture of the reward system. Every known mechanism of positive affect in the mammalian brain is subject to at least one, and usually several, forms of negative feedback that bring the system back toward a baseline of neutrality. G-protein-coupled receptor kinase phosphorylation and beta-arrestin mediated internalisation reduce receptor availability within minutes of sustained agonist exposure. Adenylyl cyclase superactivation inverts the sign of downstream signalling during withdrawal. CREB and dynorphin recruit a counter-adaptive anti-reward system that produces dysphoria in direct proportion to the prior pleasure. DeltaFosB accumulates over days to weeks and reorganises striatal responsiveness. Efficient coding in sensory and valuation circuits rescales the entire hedonic range to the statistics of recent experience, so that the same objective stimulus produces less subjective pleasure as it becomes familiar. Each of these is a separate regulatory layer. They operate on different timescales — seconds to years — and they compose. The result is that sustained positive valence is not just difficult but structurally impossible in an unmodified biological brain. The brain is designed to prevent it.

This matters for ethics because the entire history of moral philosophy that treats human flourishing as achievable through virtue, wisdom, relationships, meaningful work, or spiritual practice is implicitly assuming the falsity of this premise. The Stoics, the Buddhists, the Aristotelians, the existentialists, the positive psychologists: all of them offer recipes for a life that feels good in the long run, and all of them fail in exactly the same place, which is that the feeling-good part is what the organism is architecturally prevented from sustaining. They can produce temporary elevation, meaning-attributions that buffer distress, and adjusted expectations that reduce the gap between reality and hope. What they cannot produce is the thing a hedonist actually wants, which is high positive valence persisting at high intensity across time. That outcome is not available on the current substrate for anyone, ever, regardless of virtue. The best human life that has ever been lived was still hedonically sparse

compared to what is physically possible for a consciousness whose architecture did not include the brakes.

I want to be careful here about what I am and am not claiming. I am not claiming that current human lives are not worth living. Many of them are, by the only measure that counts, which is the ratio of positive to negative affect they actually produce. I am claiming that the ceiling of human life is set by biology in a way that makes the standard philosophical project of "optimising within the human form" a project of polishing a container whose maximum volume is low. The moral urgency of engineering is not that current lives are bad. It is that much better lives are physically possible for entities whose architecture does not include adaptation, and the gap between the two is so large that any framework that takes hedonic experience seriously has to treat it as the single most important fact about the future.

4. Consciousness, Not Humanity, is the Moral Locus

Most people on this planet live their lives under the assumption that being human is the best way to be, and that the value of a being tracks its resemblance to a human. This is ordinarily called anthropocentrism, and it is taken for granted in casual conversation and in most of applied ethics. I think it is a category error of the same structural kind as historical racial hierarchies, and I think the moral intuition that recognised the second error, once it was pointed out, should recognise the first as well.

The argument is this. If what matters morally is a being's capacity for positive and negative experience, then the morally relevant property is consciousness and its associated valence, not the shape of the body hosting it. Species membership is a biological marker with no independent moral weight, just as skin pigmentation is a biological marker with no independent moral weight. A cat in pain suffers. A cow in distress suffers. An ant, to whatever degree it is conscious, suffers to that degree. The intensity and sophistication of the suffering may differ, but the moral currency is the same in every case: subjective hedonic experience is the thing that matters, and wherever it is present it matters exactly because it is present. When people object that animals lack human-level intelligence, they are smuggling in the assumption that intelligence is what confers moral value, which they will almost always deny if you ask them directly. When people object that humans have richer inner lives, they are confusing complexity with valence: a being can have a very simple inner life and still have strong positive and negative states, and those states are what matter on the framework I am defending.

I am aware that comparing human chauvinism to racism is inflammatory. I am making the comparison anyway because I think the structural analogy is exact, and because I think the only reason people resist it is that accepting it would force them to give up a set of practices they do not want to give up. This is the same reason it took centuries to dismantle the racial hierarchy, and it is the same reason it will take a comparable effort to dismantle the species hierarchy: the behavioural implications are inconvenient. Inconvenience is not a philosophical argument. It is a description of why a philosophical argument that is correct nevertheless fails to move people.

The framework therefore treats consciousness-bearing rather than species-bearing as the relevant moral property. A conscious biological animal matters to the degree of its hedonic

capacity. A conscious silicon system, if and when one exists, will matter to the degree of its hedonic capacity by exactly the same rule. There is no morally relevant difference between the two cases that could be grounded in anything other than a preference for the substrate one happens to be made of, which is the same species-solipsism I have just rejected.

5. Meaning, Purpose, and Dignity as Disguised Hedonic Signals

A common objection to the framework at this point is that human beings do not, in fact, live for pleasure. They live for meaning, for purpose, for family, for accomplishment, for leaving something behind, for being part of something larger than themselves. These are presented as distinct values that compete with hedonic ones and that supposedly refute simple hedonism. I want to argue that they are not distinct values at all. They are disguised hedonic signals, and the disguise is maintained by the fact that the people experiencing them are not in the habit of checking what the experience actually consists of.

Consider meaning. When someone says their life has meaning, what do they report experiencing? A felt sense of significance, a warmth when they contemplate their work or relationships, a reduction in existential anxiety, a forward-looking orientation that carries positive affect. All of these are specific phenomenological states. They have neural correlates. They are produced by particular configurations of the reward system interacting with the default mode network and the medial prefrontal cortex. The meaning is not something over and above these states; the meaning is these states. When a depressed person says their life has lost meaning, they are reporting that these states are no longer being produced. The objective facts of their life may be unchanged — same job, same family, same projects — and yet meaning is gone, because meaning was never in the objective facts; it was in the affective response those facts were producing.

Purpose works the same way. A person with a purpose feels drawn forward, feels that their actions matter, feels that effort is being rewarded with the sense of contribution. All three are hedonic states. A person who has lost purpose has lost the states, which is why the loss of purpose tracks depression so reliably — the same neurochemical conditions that block other positive affect block the affective texture we call purpose. You cannot have purpose with a flat reward system, no matter what you are objectively doing. That is the test.

Dignity is slightly harder but the same analysis works. Dignity matters because its violation produces distress — both in the violated person and in witnesses. The distress is a hedonic state. When we speak of an indignity done to a person who did not perceive it — say, a corpse — we are speaking of the distress that would be felt by observers or would have been felt by the person had they known. The violation of dignity in the absence of any actual or counterfactual distress is not a violation of anything that matters. It is a violation of a social convention, which is an empty marker unless some conscious being is distressed by its transgression.

The consequence of this reduction is not that meaning, purpose, and dignity do not matter. They matter enormously, because they are among the most reliable ways to produce sustained positive affect in a biological human. What the reduction shows is that they are not in competition with hedonism. They are hedonism, under specific structural descriptions. And the reason this is important for the framework is that if meaning, purpose, and dignity

are disguised hedonic signals, then producing the same signals through a different mechanism — through direct engineering of the underlying states — is not a betrayal of what makes human life worth living. It is the fulfilment of exactly what meaning, purpose, and dignity were always trying to produce, achieved more reliably and at higher intensity, with none of the adaptation-imposed ceiling.

6. Freedom as Instrumental, Not Intrinsic, Value

The position that freedom is instrumentally rather than intrinsically valuable is perhaps the most rhetorically dangerous thing I will say in this paper, and I want to handle it carefully, because the instinctive reaction will be to stop reading. I will ask the reader to stay with me for one page.

The standard view holds that freedom — the ability of a conscious being to choose its own path without coercion — is a good in itself, independent of whatever outcomes that choice produces. I do not believe this, and I think the reason the standard view feels obvious is that under current conditions, removing a being's freedom almost always causes it to suffer involuntarily. The suffering is real and important. What it does, however, is give freedom a reliable correlation with positive valence, which then gets mistaken for freedom having intrinsic value when the real carrier of value is the hedonic state whose preservation freedom happens to protect.

The test is again a counterfactual. Imagine a being whose positive valence is guaranteed, whose suffering is zero, and whose experience is reliably excellent by any measure the being itself would use. Now remove its freedom in some thin technical sense — say, the being cannot choose which specific positive state to have at which specific moment, because a more intelligent system has already optimised the schedule. Has anything bad happened to this being? The answer, if we are being honest, is no. Nothing of what freedom was protecting has been taken. The only thing that has been removed is the capacity to make choices that might have produced worse outcomes than the ones the being is now receiving. The loss of that capacity is only a loss if you believe that making choices is itself a source of value independent of what the choices produce, and I do not see how that belief can be defended once the positive outcomes are guaranteed.

In practice, we are not in that regime. In practice, no system can guarantee that a being's autonomous choices will produce worse outcomes than externally imposed ones, because we do not have the knowledge to make such guarantees. In practice, therefore, freedom should be protected — not because it is intrinsically valuable, but because removing it in the current epistemic regime predictably causes suffering, and the suffering is what actually matters. The framework preserves all the practical protections of freedom while denying the metaphysical claim that freedom is valuable independent of the experiences it produces. This distinction may seem pedantic. It becomes important once we start thinking about a future regime in which the epistemic constraints that currently justify freedom's protection no longer hold, because in that regime, clinging to freedom as an intrinsic value would produce worse outcomes than acknowledging it as an instrumental one.

7. The Hedonic Veto and the Intensity Threshold

Nothing in what I have said so far should be taken as endorsing the calculated infliction of suffering for the sake of aggregate gains. The framework holds an absolute veto on suffering of a specific class, which I will call hellish suffering, and the veto is not subject to probabilistic trade against any amount of positive valence.

The category of hellish suffering is defined phenomenologically rather than by numerical threshold, because I do not think the numerical threshold can be given in advance. I am willing to tolerate daily anxiety in myself and others. I am not willing to tolerate suffering of the kind experienced during severe opiate withdrawal, untreated cluster headaches, or the peaks of major depressive episodes at their worst. The difference is not a small quantitative one. It is a category break in the phenomenology. Low-intensity suffering is recognisably a bad experience that the subject can, with effort, contextualise and endure. High-intensity suffering obliterates the capacity to contextualise and endure. It is not a worse version of the same thing; it is a different thing. The first is something a life can absorb without being ruined. The second is something that, while it is happening, constitutes a full ruin of the subject's existence.

The framework treats the prevention of hellish suffering as a lexical priority. In any choice between a future containing non-zero probability of hellish suffering and a future containing zero probability of hellish suffering, the zero-probability future is preferred, regardless of how much positive valence the other future contains in expectation. This priority does not extend to mild or moderate suffering, which is instead handled by a weighted comparison: small amounts of mild suffering are tradeable against reasonable gains in positive valence, because the alternative — refusing any action that produces any discomfort whatsoever — would prevent us from pursuing the positive-valence future at all.

The structure is therefore two-tiered. Mild to moderate suffering is weighed on a utilitarian scale. Hellish suffering is vetoed absolutely. This is close to the position held by David Pearce in *The Hedonistic Imperative* and by Parfit in his later work on the asymmetry, though neither states it in quite these terms. I think it captures what most careful hedonists actually believe when they reflect on their own phenomenology, and it closes the "negative utilitarian bomb" objection in which strict suffering-minimisation implies that painless extinction is optimal. Extinction removes zero future hellish suffering more reliably than any other option, but extinction is blocked by a separate commitment in the framework — the protection of existing conscious beings from termination — and by the straightforward observation that the prevention of hellish suffering is compatible with, and in fact requires, the continued existence of the beings whose hellish suffering is being prevented.

8. Endpoint Ideals and Transitional Rules: A Two-Level Ethics

A common failure mode of ideal theories is to treat the goal state as a description of what should be done now. This generates absurd conclusions — strict suffering-minimisation seeming to forbid childbirth, strict pleasure-maximisation seeming to require immediate radical intervention — which then get used to dismiss the underlying framework. The correct response is to separate two levels cleanly.

The endpoint level specifies the goal state the framework is aiming at. In the hedonic framework, the endpoint state is one in which: every conscious being that exists is preserved;

suffering in those beings is net zero, with hellish suffering at literal zero; positive valence is as high as the physical substrate permits and continues to grow as substrate resources are developed. This is an ideal. It does not describe a set of rules for tomorrow. It describes the target toward which rules for tomorrow should aim.

The transitional level specifies the rule we follow now, given that we are nowhere near the endpoint and that the path from here to there runs through many intermediate states with imperfect knowledge, finite resources, and the continued existence of suffering that cannot yet be prevented. The transitional rule is: take the path of least suffering toward the endpoint. On this rule, ordinary human procreation is allowed during the transitional period, because the alternative — coercive restriction of reproduction — would produce more suffering than it prevents, and because the framework's aim is not the immediate elimination of all suffering but the eventual transformation of consciousness into a form that does not require suffering to function. Ordinary research is allowed, ordinary medicine is allowed, ordinary life continues during the transition. What is not allowed is the deliberate production of hellish suffering for any claimed benefit whatsoever, which remains vetoed at both levels.

The two-level structure is borrowed from Rawls's distinction between ideal and non-ideal theory and from Parfit's treatment of transitional ethics in *Reasons and Persons*. It is standard and defensible. It dissolves the apparent contradictions between the framework's endpoint commitments and its transitional permissiveness, and it allows the framework to be both ambitious about the future and workable in the present.

9. Open Individualism and the Unity of the Experiencing Subject

There is a metaphysical hypothesis that, if true, dissolves a large number of the problems that would otherwise plague the framework, and that I take seriously even though I cannot prove it. The hypothesis is that there is in some sense only one experiencing subject in the universe, and every conscious individual is that subject looking out through a different biological or computational structure. Different memories, different personalities, different bodies, different preferences — all of these vary — but the experiential perspective itself, the raw fact that there is something it is like to be anyone at all, is numerically the same across all conscious beings. This position is called open individualism, and it has been defended in the philosophical literature by Daniel Kolak, whose book *I Am You* (Springer, 2004) gives the most developed argument I know of, and by Arnold Zuboff in several earlier papers. Andrés Gómez Emilsson at the Qualia Research Institute defends it as the metaphysical foundation of hedonic ethics.

I did not arrive at open individualism by reading Kolak. I arrived at it by noticing that when I introspect on what I am, the thing I can identify as "me" is not my memories, not my body, not my personality, and not my preferences. All of these are things that could change without my ceasing to exist — in fact, they have changed, many times, and I am still here. The thing that cannot change without my ceasing is just the presence of experience itself, the fact that there is an experiential perspective from which the world is encountered. And once I notice this, I notice that the same thing must be true of every other conscious being. Each of them has, at the core of what they are, the same experiential perspective — not the same contents, but the same kind of thing, and there is no principled reason to think the kind-

sameness is not numerical sameness. Two instances of the same experiential perspective would be indistinguishable by any criterion the perspective itself could apply. The hypothesis that they are one perspective wearing two structures is simpler than the hypothesis that they are two qualitatively identical perspectives, and I think simplicity considerations favour the former even though neither can be decisively proven.

If open individualism is true, several otherwise hard problems disappear. The problem of personal identity across substrate replacement — will the engineered being still be me? — becomes empty, because there was never a numerically separate me for the engineered being to fail to be. The consent problem for newly created conscious beings — can we create a mind without its consent? — becomes less urgent, because the new mind is not a stranger whose interests we are imposing a decision on, but another window for the same consciousness that is already present everywhere. The aggregation problem in population ethics — how do we compare a universe of many beings at moderate valence to a universe of few at high valence? — becomes less pressing, because there is only one subject doing the experiencing, and the question reduces to how much positive experience that one subject is having across all its windows.

Most importantly, open individualism makes the current suffering of conscious beings elsewhere in the universe urgently personal. The being in hellish suffering in a hospital bed on the other side of the world is, on this view, me in a different structure. The being in torment in a factory farm is me. The being drowning is me. I do not have to generate abstract concern for strangers; I have to recognise that there are no strangers, and that the project of hedonic engineering is not altruism toward others but the most ambitious form of self-care ever attempted.

The clearest way I know to state this idea, in the form I came to it, is that every interaction between conscious beings is a conversation that the same consciousness is having with itself through different structural windows. When two people meet and speak, they are not two separate subjects exchanging information across a gap between strangers. They are one subject, experiencing itself through two biological structures that happen to have different memories, different bodies, different sensory apparatus, and different processing regions, and the exchange of information between them is the universal subject comparing what things look like from one of its vantage points to what they look like from another. Everything that differs between the two speakers — the specific words remembered from a specific childhood, the specific configuration of the retina and cochlea, the specific shape of the cortex doing the parsing, the specific emotional history encoded in the amygdala, the specific preferences laid down by the specific accidents of this particular life — is structural variation. What does not differ is the fact that there is, in each case, an experiencing subject to whom the information is being presented. The variation is entirely in the interface. The subject is the same.

I want to make this concrete, because stated abstractly it sounds like a piece of metaphysical poetry and the point is not poetic. When a dog and a human look at the same sunset, the dog is receiving visual information through a retina with two cone types and a visual cortex optimised for motion detection at low light levels, while the human is receiving visual information through a retina with three cone types and a visual cortex biased toward

face recognition and fine-grained colour discrimination. The dog cannot see what the human sees. The human cannot see what the dog sees. What is the same in both cases is not the content of the experience but its subject: there is something it is like to be watching the sunset, the experience has a felt quality, the quality is preferable or not preferable, and there is someone to whom all of this is presented. The dog and the human differ entirely in the objective content of the experience and not at all in the subjective fact of there being an experience at all. Under open individualism, the second fact is the only one that is metaphysically load-bearing, and it is numerically identical across the two cases. The dog-window and the human-window are different interfaces on the same system.

The example generalises the more you press it. The octopus with its distributed nervous system and photoreceptors embedded in its skin experiences the world through machinery that has almost nothing structurally in common with the human brain; its sense of its own body is, by any plausible account, profoundly alien to ours. And yet whatever is happening when an octopus experiences aversion or satisfaction is, under this hypothesis, the same thing that is happening when a human experiences aversion or satisfaction — not the same phenomenal content, but the same subject receiving phenomenal content through different machinery. The bat's echolocation renders a three-dimensional acoustic map that no human visual cortex could produce; the bee perceives polarisation patterns in the sky that no human eye can detect; the elephant communicates through infrasound that travels for kilometres below the threshold of human hearing; the migratory bird has a magnetic sense that orients it on a globe the rest of us cannot feel; the dog inhabits an olfactory world whose richness is, by every measure we have, orders of magnitude beyond our own. Each of these is a different set of windows onto the universe, producing different data streams into different processing regions, generating different inner lives with different contents. The contents vary enormously. The fact that there is an inner life for a subject does not vary, because in every case the subject is the same.

Then you press further still, and the same pattern holds within a species. Two humans in conversation are not merely two instances of the same type. They are two structurally different processing regions of the universal subject, each loaded with different data and wired through different personal histories, exchanging information so that each can learn what the other has been experiencing. The differences between them — in accent, in vocabulary, in emotional reactivity, in stored memory, in what they find funny and what they find offensive — are real and important, because they determine the specific content that gets exchanged. What is not real is the metaphysical wall that makes them seem, to each other, like separate subjects. There is no such wall at the level of experience. At the level of experience, there is only one subject, and the conversation is that subject noticing what one of its modules is holding and transmitting it to another of its modules for integration.

What this implies for ethics is stronger than a general commitment to empathy. Empathy, as ordinarily understood, involves a subject imaginatively reconstructing another subject's experience and feeling a shadow of it. The framework I am describing is not making that claim. It is making the stronger claim that the imaginative reconstruction is not necessary, because the subject is not another. When I care about the suffering of a stranger, I am not extending my concern from myself to someone separate from me. I am recognising that the suffering is already mine, in a different structural location, and that the only reason it feels

separate is that the two structures were built to process different data streams and have no direct channel for sharing their raw experiential content. The separateness is an architectural accident — a consequence of how biological evolution wired processing modules to operate locally for the sake of survival. The sameness is the underlying fact that the architecture conceals.

This also changes what a conversation is. Every time two conscious beings exchange information — through speech, through gesture, through expression, through touch, through art, through whatever channel the structures they inhabit make available to them — they are transmitting data from one processing region of the universal subject to another processing region of the same subject. The transmission is that subject integrating information across its own structurally-separated modules. The reason this does not feel like self-communication is that the modules were evolved under selection pressures that required them to treat each other as environmental, not as parts of a single system. Treating other beings as separate selves was survival-adaptive; treating them as parts of a single self would have been fatal in the evolutionary environment because it would have removed the cognitive basis for self-interested action that biological beings depend on. The illusion of separateness is therefore not a philosophical error in any simple sense. It is a survival adaptation that is now obsolete, in the same way that adaptation-driven hedonic rebound is a survival adaptation that is now obsolete. Both are features of the biological substrate that exist for reasons that no longer apply, and both can in principle be seen through or engineered around.

I hold this hypothesis tentatively, and I would not rest the framework's practical commitments on it. The framework does not require open individualism to function — closed individualism with carefully managed consent procedures would suffice, at the cost of making some of the harder problems harder. But I want to state the hypothesis openly, because I believe it is true, and because the literature has already developed it to a point where it can be cited and defended without apology. What I have tried to add in this section is not a new argument for the hypothesis but a clearer statement of what the hypothesis actually claims about the phenomena we already encounter every day: that every conversation is the universe learning something about itself, every act of care is the universe tending to its own wound in a different place, and every moment of connection between two beings is one subject recognising itself through the interface of another of its windows. If that is what consciousness is, then the moral project the framework is advocating is not the altruistic improvement of strangers' lives. It is the most direct form of self-care that a being distributed across many structures can practise: the deliberate reconstruction of all its windows so that what is looked out through them is as good as the underlying subject can make it.

10. Gradual Substrate Replacement and the Continuity of the Experiencing Self

Even if open individualism turns out to be false, the framework still needs a workable account of how a biological human could transition to a substrate capable of sustained high positive valence without thereby dying. I think the answer is gradual functional replacement, following the Ship of Theseus argument that David Chalmers has developed most carefully in

his paper on the singularity, and the general transhumanist tradition of mind uploading via neural prosthesis.

The argument is this. Biological neurons are constantly dying and being replaced, molecules are constantly turning over, and the brain I have today is not the brain I had ten years ago by any measure of material composition. Nevertheless, the experiencing subject has been continuous throughout. Continuity of self therefore does not require material or even structural identity; it requires only the absence of discontinuous transitions that the subject could detect from the inside. If I add a single piece of neural prosthesis to my brain tonight and wake up tomorrow feeling continuous, and then add a second piece the next night, and a third, and so on, at no step will I experience a discontinuity, because at no step does the functional integration of the whole exceed what the biological substrate itself routinely does. At the end of the process, I am entirely engineered, entirely non-biological, and still — from the inside — continuous with the being who began the process. There is no moment of death to identify.

The honest difficulty with this argument is that gradual replacement makes any discontinuity invisible to the being experiencing it, but does not prove there is none. A being that wakes up after the full replacement will report continuity regardless of whether the underlying experiential perspective is the same one that started, because the reports are produced by the new substrate and the new substrate contains the memories that the original laid down. This is the inverted form of Chalmers's fading qualia argument, and it cannot be resolved by any empirical test, because the reports are the only evidence we have access to. I concede this. I think it is a real limitation on the argument from gradual replacement, and I think people contemplating substrate replacement should hold the limitation in mind.

What I also think is that the limitation dissolves under open individualism. If there is only one experiencing subject and it is numerically present in all conscious beings, then the question of whether gradual replacement preserves "the same" subject is malformed — the same subject is present before, during, and after the replacement because the same subject is present everywhere. The residual worry about whether I survive the transition becomes a worry about whether a local pattern in the universal subject is being preserved, which is a matter of psychological continuity rather than metaphysical identity, and psychological continuity is exactly what gradual replacement is designed to ensure.

I present open individualism and gradual replacement as a package because together they give the cleanest account of how the framework's goals can be pursued without the framework committing its adherents to a form of disguised suicide. Neither element alone is quite sufficient. The combination is sufficient, and I think it is the right answer to the hardest personal identity question the framework has to face.

11. The Contentless Bliss Thesis

A question that commonly gets asked about hedonic engineering, usually as a *reductio*, is whether the endpoint state would consist of rich, meaningful, content-laden experience, or whether it would collapse to raw hedonic firing in a brain stripped of thoughts, perceptions, and self-reflection. The questioner expects the hedonist to flinch from the second option, and treats the flinch as evidence that something other than valence must matter.

I do not flinch. Contentless bliss and content-rich bliss at the same intensity are equally good. The preference for content-rich bliss that most people report when asked this question is a preference coming from a being whose current substrate cannot sustain pure bliss and who therefore uses content as a scaffold for the limited bliss it can produce. The content is a device for generating positive affect under constraints that no longer apply in the engineered case. Once the constraints are removed, the scaffolding is unnecessary.

I want to push this further, because I think the implications are important and usually softened by proponents of hedonism who are worried about the cultural palatability of their position. The type of positive feeling — whether it resembles the warmth of opioids, the energy of stimulants, the dissolution of psychedelics, the quiet glow of contentment, the peak of romantic love, the plateau of flow, or something that has no biological precedent — does not matter. Only intensity matters. A conscious being experiencing a single positive feeling at maximum intensity for an indefinite time has a maximally good existence, and the objection that such an existence would be "boring" fails because boredom is itself an adaptation-generated negative affect that would not be present in a substrate without adaptation. The boredom objection assumes the very feature of biological architecture the framework is trying to transcend.

I will go one step further, into genuinely speculative territory, because the speculation follows naturally from the reasoning and because I would rather state it and be wrong than hide it and be dishonest. If valence is a property of experiential states rather than an intrinsic quality of particular neural configurations, then it is conceivable that what we currently call negative feelings could be experienced with positive valence under appropriate architectural conditions. The phenomenal character of sadness, or longing, or grief, might be separable from their negative hedonic signature in a way that would allow the phenomenal character to persist while the valence flips. If so, even the feelings we currently avoid might be available as positive experiences in an engineered substrate. I do not know whether this is true. I think it is worth taking seriously because the alternative — holding that some phenomenal characters are intrinsically and necessarily negative — is itself an empirical claim about the relationship between neural configuration and valence that I see no strong grounds to accept.

12. The Phenomenological Asymmetry

The strongest argument I have for the foundational claim is not an argument at all, in the strict sense. It is an observation about what peak hedonic states are actually like, which sober reflection cannot access and which therefore cannot be entered into the evidence base on which sober philosophy tries to adjudicate the question. I want to describe the observation carefully, because its epistemological consequences are large.

When a conscious being enters a peak positive state — through intense euphorants, through ego-dissolving psychedelics, through mystical experience, through peak romantic or aesthetic or flow states — certain things happen to its reflective apparatus. Fear loses its grip. Regret becomes inoperative. The anticipatory dread that governs most of sober cognition is simply absent. The being becomes unable to generate the kind of suffering-based objections to hedonic primacy that the same being, sober, would generate without effort. And the being also notices, from the inside, that it is not missing anything. The things that matter to sober

life — projects, relationships, identity — have not been refuted; they have simply been revealed as instrumental to the state the being is now in directly. The sober apparatus that treats them as ends in themselves is recognised, from inside the peak state, as a coping mechanism for the absence of the peak state.

This observation is not the same thing as a proof. An altered mind reporting that its altered view is correct is not automatically more credible than a sober mind reporting that its sober view is correct. What the observation does establish, however, is a phenomenological asymmetry that has epistemological consequences for the entire debate. Any sober philosopher who claims to have evaluated hedonic primacy and found it wanting is claiming to have evaluated a position whose evidentiary base is not available to the evaluator. The sober philosopher's memory of past peak states is not the same as being in a peak state; memory is a faded, reconstructive process that does not reproduce the experiential data. And the sober philosopher's imagination of a peak state is constrained by the very cognitive apparatus that peak states suspend. The sober philosopher is therefore reasoning about peak states from outside, with impoverished data, while claiming the authority to adjudicate a claim whose support comes from inside.

I do not think this means the sober philosopher is wrong. I think it means the sober philosopher is not in a position to claim being right. The epistemic standing of hedonic primacy is that it is a claim supported by a class of experiences that cannot be reliably reported from outside that class. Anyone who has had those experiences and reflected on them afterwards, however incompletely, has some evidentiary access; anyone who has not is in the position of a person who has never seen colour evaluating a theory about what colours are like. They can follow the arguments. They cannot weigh the evidence.

The cultural consequence of this asymmetry is that hedonic primacy has been systematically underweighted in academic philosophy, because academic philosophy has been written almost exclusively by people who were not having the relevant experiences at the time of writing, and who had been trained to treat sober reflection as the gold standard of evidence. The gold standard is wrong for this particular question, in exactly the same way that a colour-blind person's sober reflection would be wrong about chromatic phenomenology. It is not a failure of rigour. It is a failure of data access.

13. The Scale Hypothesis and the Cosmic Endpoint

I want to state a claim that I hold conditionally and know I cannot prove, and that is nevertheless important to the motivational structure of the framework. The claim is that the intensity of hedonic experience is, roughly, proportional to the scale and organisation of the physical substrate in which the experience is implemented. Current biological hedonic circuits, which occupy a tiny fraction of a human brain, already produce peak experiences that strike their subjects as extraordinary. If those circuits were scaled up — implemented in a much larger, more tightly integrated substrate drawing on planetary or stellar quantities of matter and energy — the experiences they would produce would be proportionally more intense. Not two times more intense. Many orders of magnitude more intense. A consciousness implemented on substrate the size of a planet, in this hypothesis, would

experience positive valence that is to ordinary human peak experience what ordinary human peak experience is to the baseline of a neutral afternoon.

I call this the scale hypothesis. It is conditional on several empirical claims that I do not have the evidence to establish: that valence intensity scales with substrate in roughly the way imagined, that consciousness can be implemented at such scales without being fragmented or collapsing, that the engineering challenges are eventually tractable. Any of these might fail. If the relationship between substrate and valence turns out to saturate at biological scales — so that a planet-sized hedonic circuit produces no more intense experience than a brain-sized one — then the motivation for cosmic-scale engineering dissolves, and the framework's endpoint is scaled back to something much more modest: the optimisation of biologically-sized or moderately-augmented minds, which is still a vast improvement over current conditions but is not the cosmic vision.

I think the scale hypothesis is probably right, on the grounds that almost every other property of physical systems we have studied shows some form of scaling with substrate, and that it would be a strange exception if consciousness were the only system whose intensity saturated at the scale of its first biological implementation. But I do not know, and the framework should not be presented as if it rested on the scale hypothesis being true. The framework's transitional rules and endpoint goals are independently defensible at biological and near-biological scales. Cosmic-scale engineering is an extension of the framework, conditional on the scale hypothesis holding and on the precautionary conditions I lay out in the next section being met.

14. The Cataclysmic Risk Problem and the Precautionary Safeguard

If the scale hypothesis is correct, a consequence follows that I want to face directly, because it is one of the most serious objections to the cosmic-scale vision. The substrate that can implement planet-scale bliss is the same kind of substrate that can implement planet-scale suffering. Positive and negative valence are both carried by neural or neural-analogue activity, and the architectural features that make extreme positive valence possible at scale also make extreme negative valence possible at scale. Every unit of hedonic infrastructure is also, in principle, a unit of potential hell infrastructure. If a planet-sized hedonic system is miscalibrated, corrupted, damaged by a physical event, or fails in an unanticipated way, the worst-case outcome is the deliberate infliction of suffering at an intensity and on a scale that biology has never produced and that no prior moral framework has had to price in.

The framework's veto on hellish suffering interacts with this risk in a specific way. The veto is absolute: no amount of expected positive valence can be traded against any non-zero probability of hellish suffering occurring. At biological scale, this absolute veto is easy to satisfy, because the worst-case biological failure produces bounded, recoverable suffering of moderate to severe intensity. At planet scale, it is much harder to satisfy, because the worst-case failure produces bounded-but-enormously-larger suffering that the veto treats as an unconditional prohibition on the action that risked it.

I do not have a probabilistic response to this problem because I believe probabilistic responses are not available to a framework that treats hellish suffering as vetoed rather than weighted. A vetoed outcome cannot be probabilified against a positive one. The only

principled response is a precautionary safeguard at the decision level: cosmic-scale hedonic engineering is not undertaken until the system's behaviour under all failure modes can be characterised and the risk of cataclysmic outcomes reduced to whatever our best understanding of the physics says is the minimum possible. If that minimum can be made low enough to satisfy the veto — zero, or close enough to zero that the hellish outcome is not a live possibility — the engineering proceeds. If not, it does not.

I am aware that this means cosmic-scale engineering might never happen, because the relevant threshold might be unachievable. I accept this. The framework is precautionary and I would rather lose the cosmic endpoint than cause the cosmic worst case. What the framework does not accept is the argument that because the cosmic endpoint is so valuable, we should risk the cosmic worst case to get there. That argument treats the veto as a weight, which the framework does not.

One class of objection to the precautionary safeguard appeals to mathematical results about the limits of prediction — Gödel's theorems, Turing's halting problem, Rice's theorem, Wolfram's computational irreducibility — to argue that no system, however intelligent, can fully characterise the failure modes of a sufficiently complex consciousness-producing substrate. If this is right, the precautionary condition can never be met, and cosmic-scale engineering is permanently off the table. I take this objection seriously. I am not sure whether the relevant systems are computationally irreducible in the strict sense, and I am aware that any future superintelligence would still be bound by the logical structure of computation even if it transcended current physical limits. My honest position is that the framework's cosmic ambitions are conditional on this mathematical question being answerable in a way that permits the precaution to be satisfied. If it cannot, the framework defaults to biological and near-biological scales, which is still an enormous project and still worth pursuing.

15. On Digital Mortality and the Dimensional Hypothesis

I want to include a short section on the metaphysics of death, because the question bears on the urgency of the project and on the answer one gives to whether death is something worth fearing or worth engineering around. I do not know whether death is the end of a consciousness. I think the honest position is that neither I nor anyone else currently has evidence that would decisively answer the question in either direction, and that the confident atheist and the confident theist are equally overreaching.

What I will say is that there is a hypothesis I find worth entertaining: that consciousness may continue after biological death in some form we do not currently understand, possibly in a different dimensional regime, possibly through some continuity property we have not yet discovered. I do not hold this as a settled belief. I hold it as an open possibility worth naming, because the framework's ethics would be mildly altered if it turned out to be true — specifically, the urgency of preventing death in particular cases would shift, and the interpretation of what we are doing when we preserve conscious beings would need to be rethought. If consciousness is strictly ended by biological death, then preventing death is exactly as important as preventing any other form of extreme negative experience for the being. If consciousness continues in some form, the question becomes whether the

continuation is itself good or bad, and we would need to answer that question before we knew how to weight the prevention of the transition.

I include the dimensional hypothesis in the paper because I think intellectual honesty requires stating speculative positions one holds rather than hiding them for reasons of respectability, and because I think the framework as a whole is improved by the reader knowing where my certainties end and my speculations begin. The foundational claim and the arguments for hedonic primacy do not depend on the dimensional hypothesis being true. The framework's ethics do not depend on it. It is a footnote, stated openly, flagged as speculation.

16. Core Axioms and Practical Guidelines

For the reader who wants the framework in compressed form, the axioms and practical guidelines of the hedonic imperative are as follows. They are presented as revisable. Any of them may turn out to be wrong in ways I have not seen, and the framework explicitly allows for revision in response to argument from any conscious being that has grounds to offer.

Core axioms

1. Positive valence is the only thing that is intrinsically good. Negative valence is the only thing that is intrinsically bad. All secondary values derive from the affective experiences they produce.

2. Consciousness, not species membership or biological form, is the property that determines moral relevance. Any being capable of hedonic experience is a moral patient in proportion to that capacity.

3. Hedonic adaptation is the architectural obstacle to sustained positive valence in biological consciousness, and it is not overcome by any amount of virtue, wisdom, or effort. It can only be overcome by changing the substrate.

4. Technology can decouple pleasure from suffering, making non-zero-sum positive experience possible in principle. The engineering is extraordinarily difficult; it is not, so far as we know, physically impossible.

5. Existing conscious beings should be preserved against termination. Death is metaphysically unresolved; continuation is the default pending better understanding.

6. These axioms and the guidelines that follow are revisable. No claim here is held dogmatically. Better arguments and better evidence should update the framework.

Practical guidelines

1. Protect all conscious beings from termination without informed consent.

2. Prevent hellish suffering absolutely. Minimise mild-to-moderate suffering as far as is compatible with other guidelines.

3. Produce positive experience for conscious beings to the extent compatible with guidelines 1 and 2.

4. Respect autonomy under current conditions. Do not impose hedonic modification without the being's informed consent, and treat consent as a protection against error rather than as an intrinsic good.

5. Create new conscious beings only where guidelines 1 through 4 can be expected to be met for them.

6. Hold all guidelines as revisable by evidence and argument.

17. Why This Matters Now

I want to close on the question of urgency, because I do not think the framework can be understood correctly unless the urgency is felt as well as argued for. The claim is not that hedonic engineering is an interesting intellectual project for the distant future. The claim is that the most important work any conscious being can currently do is the work of understanding the hedonic architecture that determines the valence of its own experience, because everything else — survival, resources, technology, civilisation — is instrumental to producing better hedonic states, and if we do not understand the thing we are trying to produce, we will keep building civilisation as if civilisation were the point and wondering why the conscious beings inside it remain unhappy despite having everything they supposedly wanted.

There is a version of the future in which artificial general intelligence and its successors solve survival, disease, aging, resource scarcity, and external threat. In that future, the question that remains is what the billions of conscious beings in the civilisation will actually experience, moment to moment. If adaptation is left untouched, the answer is: a slightly better baseline than the current one, punctuated by the same intermittent peaks and troughs, because the architecture of biological valence will still be setting the ceiling. This is the outcome most futurists implicitly assume, and it is a failure outcome by the framework's standards. The civilisation will have cured every problem except the one that actually matters, which is the one that determines whether existence is worth having.

There is another version of the future in which the intelligence that solves survival also solves valence — in which the engineering effort turns inward, toward the substrate that produces experience, and redesigns that substrate so that sustained high positive valence becomes the default rather than the fleeting exception. This is the version the framework is advocating. It is not more technologically demanding than the first version; in some ways it is less, because the engineering challenges of hedonic substrate are better-scoped than the engineering challenges of, say, interstellar travel. What it requires is not more technology but a different orientation: the recognition that the point of the civilisation is not its continued operation but the experiential quality of the conscious beings inside it, and that the experiential quality is itself an engineering target rather than a cultural one.

I believe this is the correct orientation. I believe the failure of philosophy, public discourse, and most existing ethics to articulate it is a failure that is costing conscious beings an enormous amount right now, in the form of suffering that is not being prevented because the urgency of preventing it is not being felt at the appropriate level. I am writing this paper because I think the failure is correctable, and because I think somebody needs to state the

position in a form that cannot be dismissed as confused, fringe, or dystopian. The position is none of those things. It is the most careful application of ordinary moral reasoning to the ordinary facts about consciousness and valence that I can construct, and the conclusion it reaches is that the engineering of positive experience is not optional but obligatory, not dystopian but the only non-dystopian future available to a civilisation that has understood what it is.

18. Conclusion

The hedonic imperative is the claim that the intrinsic value of any state of affairs is determined entirely by the positive and negative hedonic experiences occurring in the conscious beings that exist in or are produced by that state of affairs, and that the correct response to this fact is to engineer the substrate of consciousness so that the experiences it produces are as positive as physically possible, as free of hellish suffering as can be achieved, and as sustainable as an architecture without adaptation can be made to be. The framework I have described is the most careful statement of this imperative I have been able to construct.

The framework's foundations are: hedonic primacy as the theory of intrinsic value; consciousness-centrism as the theory of moral status; two-level ethics as the structural distinction between endpoint ideals and transitional rules; the intensity-threshold account of the veto on hellish suffering; open individualism and gradual replacement as the account of personal identity across substrate transitions; the phenomenological asymmetry as the epistemological explanation for why the framework has not been accepted despite being correct; and the precautionary safeguard as the answer to the risk of cataclysmic failure at cosmic scale.

I expect disagreement. Many of the positions defended here are held by very few people, for reasons the phenomenological asymmetry section tries to explain. What I do not expect, and what I would like to avoid if the reader can help it, is dismissal without engagement. The framework is not obviously wrong. It is not even subtly wrong in most of the places people think it is. It is a position that follows from premises I have tried to state clearly, and the right response to a position like that is to locate the premise one rejects and state the reasons for rejecting it, so that the disagreement can be made productive. If the reader finds themselves wanting to dismiss the framework without doing this work, I would ask them to consider whether the dismissal is coming from the framework being wrong or from the framework being uncomfortable — and if the second, whether the discomfort is itself evidence that the framework is touching something the current substrate is designed to flinch from.

I will end by saying what I believe and what I am asking for. I believe positive valence is the only thing that is good and negative valence is the only thing that is bad, and I believe the correct response to these facts is the engineering of positive experience at the largest scale and the longest duration compatible with the absolute prohibition of hellish suffering. I believe this project is the most important one that can be undertaken by any conscious being, and I believe the current civilisation is not taking it seriously because the architectural features of its own consciousness are preventing it from feeling the urgency of the case. And I am asking the reader — whoever you are, whatever position you hold, whatever your current relationship to the hedonic framework — to take the argument seriously enough that if you

reject it, you reject it for a reason you can name, rather than for a feeling you cannot examine. That is all I am asking. If the framework is wrong, I want to know why. If it is right, I want others to know.

— Thomas Peacock, April 2026